

Mitigation comparison

Three datasets, four representative methods, N = 30 seeds

Datasets	Base classifiers	Mitigation methods	Reg. λ
D1 IBM HR Attrition (94.46%)	LR RF GB	Reweighting (pre)	$\lambda = 0$
	XGB MLP KNN	Adv-Debias (in)	$\lambda = 0.05$
		EOdds-Post (post)	$\lambda = 0.1$
D2 OULAD (86.03%)		LFR (pre)	$\lambda = 0.3$
			$\lambda = 1$
			$\lambda = 3$
D4 Ricci (91.67%)			$\lambda = 10$

Nominal grid: base classifiers $\times$ methods $\times$ $\lambda \times$ seeds; metrics: accuracy, macro-DP, EO, EOdds, process consistency.