


Can procedural fairness be operationalised as ML metrics, shown to be empirically distinct from statistical fairness on real HR-relevant data, and used alongside it to evaluate mitigation choices?

Composite research question



Dataset coverage

	D1 Attrition (honest)	D1 PerfRating (leaky motivator)	D2 OULAD (multi-class)	D3 ACS-Income (scale + proxy)	D4 Ricci (institutional HR)
Audit	used	used	used	used	used
Procedural	used	used	used	used	used
Mitigation	used	n/a	used	partial	used
SHAP	used	n/a	used	used	n/a