

# Procedural-fairness compute pipelines

## Process Consistency

How stable is the prediction  
under semantic noise?

Trained model  $f$   
+ test sample  $x$   
(sample\_n = 500)

Gaussian noise on numerics  
 $\sigma \in \{0.1, 0.3, 1.0, 3.0\}$

JS divergence between original  
and perturbed predictions  
Mean over rows; bootstrap CI

## Voice / Representation

How much does the model use  
modifiable features?

Trained model  $f$   
+ test sample  $x$

TreeExplainer SHAP attribution  
(or KernelExplainer)

Voice = modifiable attribution  
share over total attribution  
Enrichment adjusts for feature count

## Model-Flippability

Can the prediction be reversed by  
a sparse counterfactual?

Trained model  $f$   
+ test sample  $x$  (sample\_n = 30)

Greedy  $k_{\max} = 1$  counterfactual  
over numeric features within  $\pm 1\sigma$

Sparsity:  $1 - \# \text{changed} / |F|$   
Validity:  $x'$  flipping  $\hat{y}$

## Explanation-Actionability

Does the counterfactual flip  
a modifiable feature?

Greedy  $k_{\max} = 1$  counterfactual  
reuses model-flippability output

Is the flipped feature in  $F_{\text{mod}}$ ?

Fraction over rows; bootstrap CI